

AI Governance

Functional Scope, Dashboard Boards & Use Cases

TICVAI

Professional Edition

Table of Contents

Board 1: AI Governance & Policy Management	4
Screen 1 — AI Governance Command Center	5
Screen 2 — AI Capability Registry & Ownership	7
Screen 3 — AI Risk Classification & Assessment	9
Screen 4 — AI Autonomy Level Configuration	11
Screen 5 — AI Action & Permission Policy Builder	13
Screen 6 — AI Data Access & Usage Policy	16
Screen 7 — Environment, Tenant & Scope Governance	17
Screen 8 — AI Policy Conflict, Exception & Override Management	19
Screen 9 — AI Policy Testing & Governance Simulation	21
Screen 10 — AI Governance Policy Publication & Effective Policy Map	23
Board 2: AI Approval, Human Oversight & Intervention Management	32
Screen 1 — AI Human Oversight Command Center	33
Screen 2 — AI Approval Requirement & Routing Configuration	35
Screen 3 — AI Approval Review Workspace	37
Screen 4 — Conditional Approval & Approval Conditions	38
Screen 5 — Human Review, Challenge & AI Clarification Workspace	40
Screen 6 — Escalation, Delegation & Approval SLA Management	41
Screen 7 — Live AI Execution Oversight & Human Intervention	43
Screen 8 — Human Override & Manual Control Center	45
Screen 9 — Approval & Intervention History / Decision Timeline	46
Screen 10 — Human Oversight Workflow Simulator & Readiness Center	48
Board 3: AI Explainability, Decision Trace & Audit	59
Screen 1 — AI Explainability & Audit Command Center	60
Screen 2 — AI Decision Explorer & Search	62
Screen 3 — AI Decision Explanation Workspace	64
Screen 4 — Data, Feature & Evidence Provenance	66
Screen 5 — Candidate, Rule & Decision Path Trace	68
Screen 6 — Model, Provider & AI Runtime Trace	70
Screen 7 — Governance, Approval & Human Decision Trace	72
Screen 8 — Execution & Business Outcome Trace	74
Screen 9 — AI Audit Record & Evidence Package	76
Screen 10 — AI Trace Investigation & Replay Simulator	78
Board 4: AI Risk, Compliance & Governance Monitoring	88
Screen 1 — AI Governance Monitoring Command Center	89
Screen 2 — AI Risk Register & Risk Exposure Management	91
Screen 3 — AI Governance Control Library & Control Effectiveness	93
Screen 4 — AI Policy Compliance & Violation Monitoring	95
Screen 5 — AI Data, Privacy & Usage Compliance Monitoring	97
Screen 6 — AI Quality, Behavior & Governance Drift Monitoring	98
Screen 7 — AI Governance Alert & Detection Center	100
Screen 8 — AI Incident & Remediation Management	102
Screen 9 — AI Compliance, Assurance & Governance Reporting	3
Screen 10 — AI Governance Review, Action Plan & Continuous Improvement	3


AI Governance — Board 1: AI Governance & Policy Management

Module: AI Governance & Responsible AI Control

Board: 1 of 4

Screens: 10

Board Type: Enterprise AI Governance / Policy / Risk Classification / Autonomy / Permissions / Environment Control



Board Purpose

This board establishes the **central governance authority for every AI capability in TICVAI**.

TICVAI will eventually contain AI across many areas:

- AI Configuration Assistant
- Recommendation & Personalization
- Forecasting
- Fraud Intelligence
- Reporting & AI Analytics
- Customer/Operator Assistant
- Dynamic Pricing recommendations
- CRM/Marketing intelligence
- Enterprise Knowledge
- Future AI capabilities

Each module should **not invent its own rules about what AI can or cannot do**.

Instead:

AI Governance defines the policy. The owning TICVAI module enforces it during operation.

For example:

- **AI Configuration Assistant requests:**
“Change Adult weekend price from AED 140 to AED 150.”
- AI Governance evaluates:

Module: Pricing

Action: Modify Price

Environment: Production

Risk: High

AI Autonomy: Prepare Only

Required Approval: Commercial Owner

Result:

AI may prepare the change but cannot execute it until approval is obtained.

The core architecture is:

- **AI Request / Recommendation / Action**
↓
- **Governance Policy Resolution**
↓
- **Risk Classification**
↓
- **Permission & Autonomy Check**
↓
- **Required Controls / Approval**
↓
- **Allow / Allow with Conditions / Require Approval / Block**
↓
- **Owning TICVAI Module**

Screen 1 — AI Governance Command Center

Purpose

Provide management with one centralized view of AI governance across the entire TICVAI platform.

Header KPIs

- AI Capabilities Registered
- Active Governance Policies
- Governed AI Actions Today
- Approval-Required Actions
- Blocked AI Actions
- Policy Exceptions
- High-Risk AI Capabilities
- Policy Violations
- AI Autonomy Coverage
- Governance Warnings

AI Governance Coverage

Show all AI-enabled TICVAI areas.

AI Capability	Module	Risk Level	Autonomy	Policy	Status
Configuration Assistant	Platform	High	Level 2	GOV-012	Governed
Recommendation Engine	Recommendation	Medium	Level 3	GOV-018	Governed
Dynamic Pricing Recommendation	Pricing	High	Level 1	GOV-024	Governed
Attendance Forecast	Forecasting	Low	Level 1	GOV-031	Governed
Fraud Decisioning	Fraud	Critical	Level 2	GOV-041	Governed

Risk Distribution

Visualize:

- Low
- Medium
- High
- Critical
- **Governance Status**
 - Statuses:
 - Governed
 - Policy Missing
 - Review Required
 - Exception Active
 - Restricted
 - Suspended
 - Non-Compliant
- **Recent Governance Decisions**
 - Example:

Time	AI Capability	Requested Action	Risk	Decision
14:21	AI Configuration	Modify Price	High	Approval Required
14:18	Recommendation AI	Rank Products	Low	Allowed
14:12	Fraud AI	Block Account	Critical	Human Review
14:05	Forecasting	Generate Forecast	Low	Allowed

Actions

[Create Policy](#) | [Register AI Capability](#) | [Review Risks](#) | [View Exceptions](#) | [View Governance Activity](#)

Screen 2 — AI Capability Registry & Ownership

Purpose

Maintain a controlled registry of every AI capability operating within TICVAI.

Nothing should become an operational AI capability without being identifiable and governed.

AI Capability Record

Each capability should contain:

- • Capability ID
- • Capability Name
- • Description
- • Owning Module
- • Business Owner
- • Technical Owner
- • AI Function
- • Decision Type
- • Provider / Model Reference
- • Environment
- • Data Categories Used
- • Output Type
- • Action Capability
- • Risk Classification
- • Governance Policy
- • Status
- **AI Function Types**
- Examples:
 - • Assistant
 - • Recommendation
 - • Prediction
 - • Forecast
 - • Classification
 - • Detection
 - • Optimization
 - • Content Generation
- **Configuration**
- Decision Support
- Automated Action
- Search / Knowledge Retrieval

Example

Capability

AI Configuration Assistant

Owner

TICVAI Platform

Function

Configuration / Decision Support / Controlled Action

Inputs

Admin Intent + TICVAI Configuration

Outputs

Configuration Blueprint / Configuration Plan

Can Execute

Yes — governed

Risk

High

Capability Scope

Show which TICVAI modules the AI can interact with:

- ✓ Ticketing
- ✓ Pricing
- ✓ Capacity
- ✓ Access
- ✓ Membership
- ✓ Channels
- ✓ Media

Environment Registration

Capability may have separate governance status for:

- **Development**
- **Testing**
- **Sandbox**
- **UAT**
- **Production**
- **Lifecycle Status**
 - • Proposed
 - • Under Review
 - • Approved
 - • Active
 - • Restricted
 - • Suspended
 - • Retired

Actions**Register Capability | Edit Ownership | Assign Policy | Assess Risk | Suspend Capability****Screen 3 — AI Risk Classification & Assessment****Purpose**

Classify AI capabilities and individual AI action types according to their potential business, customer, financial, operational, security and compliance impact.

Risk should not be based only on which AI model is being used.

The more important question is:

- **What can this AI actually do?**
- **Risk Levels**
- **Low**
- Informational/advisory with limited impact.
- Examples:
 - • Explain a dashboard
 - • Summarize report
 - • Generate forecast for review
 - • Suggest configuration

Medium

Can influence customer or operational decisions.

Examples:

- • Product recommendation
- • Capacity recommendation
- • Marketing recommendation
- • Operational optimization suggestion

High

Can materially affect commercial or operational configuration.

Examples:

- • Price modification
- • Product publication
- • Capacity modification
- • Membership rule modification

Critical

Could materially affect security, access, financial controls or sensitive customer outcomes.

Examples:

- • Disable access credential
- • Fraud account restriction
- • Security-policy modification
- • High-impact bulk action
- **Risk Assessment Dimensions**
- Score/evaluate:
 - • Customer Impact
 - • Financial Impact
 - • Operational Impact
 - • Security Impact
 - • Data / Privacy Impact
 - • Compliance Impact
 - • Reversibility
 - • Scale / Blast Radius
 - • Automation Level

- Explainability Requirement
- Human Oversight Requirement

Example

Action

Modify Ticket Price

Financial Impact:

- High
- Customer Impact:
- Medium
- Reversibility:
- Medium
- Scale:
- Potentially Large
- Result:
- **HIGH RISK**
- **Risk Matrix**
- Display:
- **Impact × Likelihood / Exposure**
- with configured methodology.

Risk Override

Authorized governance users may override a calculated classification, but must provide:

- New Classification
- Reason
- Approver
- Effective Date
- Review Date
- The override is audited.

Actions

[Run Assessment](#) | [Modify Classification](#) | [View Methodology](#) | [Submit for Review](#)

Screen 4 — AI Autonomy Level Configuration

Purpose

Define exactly how independently AI is permitted to operate.

This is one of the most important screens in the governance module.

TICVAI AI Autonomy Levels

I recommend four controlled levels plus a disabled state.

Level 0 — Disabled

AI capability cannot operate.

Level 1 — Advisory

AI can:

- ✓ Analyze
 - ✓ Explain
 - ✓ Recommend
 - ✓ Simulate
- AI cannot:
- ✗ Prepare executable changes
 - ✗ Execute actions
- Example:

“I recommend increasing weekend capacity.”

Level 2 — Prepare

AI can:

- ✓ Analyze
 - ✓ Recommend
 - ✓ Build proposed change
 - ✓ Validate
 - ✓ Simulate
 - AI cannot execute the production change.
 - Example:
 - AI prepares the new pricing configuration for approval.
-

Level 3 — Execute with Approval

AI can:

- ✓ Prepare
 - ✓ Validate
 - ✓ Submit
 - ✓ Execute after required approval
- Example:

Commercial Owner approves → AI orchestration executes through Pricing API.

Level 4 — Controlled Auto-Execution

AI may execute specifically approved low-risk actions without transaction-by-transaction human approval.

This must be tightly restricted.

Example:

AI automatically updates an approved low-risk internal classification within predefined limits.

It does **not** mean unrestricted AI autonomy.

Autonomy Configuration

Configure by:

- • Tenant
- • Venue

- AI Capability
- Module
- Action Type
- User Role
- Environment
- Risk Level
- Business Unit

Example

Recommendation Ranking

Production

Risk: Low

Autonomy:

- **Level 4**
- But:
- **Recommendation Strategy Modification**
- Production
- Risk: Medium
- Autonomy:

Level 2

Maximum Autonomy

Governance policy should establish a ceiling.

Example:

- **Pricing**
- Maximum permitted autonomy:

Level 3

Therefore no local administrator can configure Pricing AI to Level 4.

Actions

Set Autonomy | Compare Environments | Configure Ceiling | View Effective Policy

Screen 5 — AI Action & Permission Policy Builder

Purpose

Define granular rules controlling what AI is allowed to read, recommend, prepare, modify or execute.

This should work alongside TICVAI's existing **RBAC/PBAC**, not replace it.

Policy Structure

Use:

- **IF**
- AI Capability
-

Module

-

Action

-

Environment

-

Risk

-

User / Role / Context

THEN

Governance Decision

Governance Decisions

- Allow
- Allow Advisory Only
- Allow Prepare Only
- Approval Required
- Multi-Approval Required
- Allow with Conditions
- Block
- Escalate

Example Policy

IF

Capability = AI Configuration Assistant

AND

Module = Pricing

AND

Action = Modify Price

AND

Environment = Production

THEN

Prepare Only

-

Commercial Owner Approval Required

Another Example

IF

Capability = Recommendation Engine

AND

Action = Rank Eligible Products

AND

Environment = Production

THEN

Allow

provided Board 1 recommendation guardrails are satisfied.

Action Types

Governance should understand actions such as:

- • Read
- • Analyze
- • Recommend
- • Generate
- • Prepare
- • Create
- • Modify
- • Delete
- • Publish
- • Activate
- • Suspend
- • Execute
- • Approve
- • Export
- • Communicate
- • Restrict
- • Block
- **Permission Intersection**
 - The effective permission should be:
- **User Permission**
 - n
- **AI Capability Permission**
 - n
- **Governance Policy**
 - n
- **Owning Module Policy**
 - For example:

A user cannot ask AI to perform an action that the user himself is not authorized to perform.

Actions

Create Policy | Test Policy | Clone | Disable | View Effective Permission

Screen 6 — AI Data Access & Usage Policy

Purpose

Control which data categories each AI capability is permitted to access and for what purpose.

This complements privacy/consent controls but does not replace TICVAI's core privacy governance.

Data Categories

Examples:

- • Public Configuration
- • Venue Configuration
- • Product Data
- • Pricing Data
- • Capacity
- • Operational Data
- • Customer Profile
- • Purchase History
- • Membership
- • Loyalty
- • Payment Metadata
- • Fraud Signals
- • Access History
- • Employee Data
- • Financial Data
- • Audit Data

• Access Matrix

AI Capability	Data Category	Read	Process	Send External	Store	Status
Recommendation AI	Basket	✓	✓	Conditional	Limited	Allowed
Recommendation AI	Payment Details	✗	✗	✗	✗	Blocked
Forecasting AI	Attendance	✓	✓	Conditional	✓	Allowed
Configuration AI	Product Config	✓	✓	Conditional	Limited	Allowed

• Purpose Limitation

- Example:
- Customer purchase history may be permitted for:
 - ✓ Recommendation Personalization
 - but not automatically for:

✗ AI Configuration Assistant

External Provider Control

If an AI provider is external, policy should determine whether particular data categories may be sent.

Example:

- **Customer PII**
- External AI Provider:

Restricted

unless approved architecture and policy permit it.

Data Minimization

Policy:

Send only the minimum information required for the AI task.

Masking / Tokenization

Where supported:

- • Mask
- • Tokenize
- • Redact
- • Aggregate
- • Anonymize / Pseudonymize where appropriate

Actions

[Configure Data Policy](#) | [Test Access](#) | [View Provider Restrictions](#) | [View Usage Audit](#)

Screen 7 — Environment, Tenant & Scope Governance

Purpose

Allow different AI governance rules across TICVAI's multi-tenant and multi-environment architecture.

A rule suitable for a development sandbox may not be acceptable in production.

Scope Hierarchy

Recommend:

- **Platform Default**



- **Tenant**



- **Business Unit**



- **Venue**



- **Module**



- **AI Capability**



- **Action**

- **Inheritance**

- Example:

- **Platform**

- Production price changes:

- Maximum autonomy = Level 3



- **Tenant A**

- inherits Level 3



- **Venue X**
- cannot raise to Level 4.
- It may tighten the policy to:
- Level 2.
- **Environment Policies**
- **Development**
- More flexible experimentation.
- **Sandbox**
- Synthetic/test data and broader AI testing.

UAT

Production-like governance with test execution.

Production

Strict controls.

Example

Action	Dev	Sandbox	UAT	Production
Generate Configuration	Allow	Allow	Allow	Allow
Prepare Change	Allow	Allow	Allow	Allow
Auto-Execute Pricing	Allow Test	Allow Test	Block	Block
Customer Data Use	Synthetic	Restricted	Restricted	Governed

Tenant-Specific Restrictions

Example:

- Tenant A:
- External AI providers prohibited for customer data.
- Tenant B:

Approved external provider permitted for specified use cases.

Effective Policy

The UI should show:

Platform Policy

-

Tenant Override

-

Venue Restriction

=

Effective Governance Policy

This is critical for troubleshooting.

Actions

Configure Scope | View Inheritance | Add Restriction | Compare Environments

Screen 8 — AI Policy Conflict, Exception & Override Management

Purpose

Identify conflicting governance rules and manage legitimate temporary exceptions without bypassing governance silently.

Conflict Types

- Policy vs Policy
- Platform vs Tenant
- Autonomy Conflict
- Permission Conflict
- Data Access Conflict
- Approval Conflict
- Environment Conflict

Example Conflict

Platform:

- Pricing maximum autonomy = Level 3
- Tenant:
- Pricing AI autonomy = Level 4
- Result:
- **Conflict — Platform Ceiling Applies**
- Effective:
- **Level 3**
- **Exception Request**
- Example:

Request

Allow AI Configuration Assistant to execute a specific bulk product update during controlled migration.

Exception Details

- Requested Capability
- Policy Being Overridden
- Requested Scope
- Reason
- Requested By
- Environment
- Start Date
- Expiry Date
- Risk
- Compensating Controls
- Required Approvers

Exception Lifecycle**Requested**→ **Risk Reviewed**→ **Approved / Rejected**→ **Active**→ **Expiring**→ **Expired****Temporary Only**

Exceptions should generally require an expiry.

Avoid:

Permanent exception with no review.

Emergency Override

For appropriately authorized operational scenarios, support controlled emergency override with:

- • Elevated permission
- • Mandatory reason
- • Limited scope
- • Time limit
- • Full audit
- • Post-action review

Actions

Resolve Conflict | Request Exception | Approve Exception | Revoke | View History

Screen 9 — AI Policy Testing & Governance Simulation

Purpose

Allow administrators to test governance policies before activating them.

This is extremely important because a policy mistake could either:

- • Allow AI too much authority, or
- • Block legitimate TICVAI operations.

- **Simulation Scenario**

- **User**

- Venue Administrator

- **AI Capability**

- AI Configuration Assistant

- **Request**

- Increase Adult Ticket Price

- **Module**

- Pricing

- **Environment**

- Production

- **Risk**
- High
- **Policy Evaluation**
- Step 1:
 - User Permission
 - ✓ Can modify pricing
- Step 2:
 - AI Capability Permission
 - ✓ Can prepare pricing change
- Step 3:
 - Autonomy
 - Level 2 — Prepare
- Step 4:
 - Production Policy
 - Approval required
- Step 5:
 - Risk Policy
 - Commercial Owner approval
- **Final Decision**
- **ALLOW PREPARATION**
- **Execution**
- **BLOCKED UNTIL APPROVAL**

Explain Why

Pricing modifications in Production are classified High Risk. AI Configuration Assistant has Level 2 autonomy for this action. Commercial Owner approval is required before execution.

Additional Test Scenarios

Scenario A

Recommendation AI ranks eligible products.

Result:

- **Allowed**
- **Scenario B**
 - Recommendation AI changes ticket price.
 - Result:
- **Blocked**
- **Scenario C**
 - Fraud AI recommends account review.
 - Result:

Allowed

Scenario D

Fraud AI permanently blocks customer account.

Result:

- **Approval Required / Restricted**, according to policy.

- **Batch Policy Testing**
- Test policy changes against:
 - Existing AI capabilities
 - Existing action types
 - Tenant configurations
 - Environments
- **Impact Preview**
- Example:
- Proposed policy change would affect:
 - **12 capabilities**
 - **4 tenants**
 - **1,840 estimated daily AI actions**
- Potential newly blocked:

220

Actions

[Run Simulation](#) | [Test Policy Set](#) | [Compare Results](#) | [View Impact](#) | [Return to Policy](#)

Screen 10 — AI Governance Policy Publication & Effective Policy Map

Purpose

Provide the final review, approval, versioning and publication layer for AI governance policies.

No material governance policy should become active without a controlled lifecycle.

Policy Package

Example:

- **TICVAI AI Governance Policy Set**
- Version:
 - **v1.4**
- Environment:
 - **Production**
- Status:
- **Pending Publication**
- **Package Summary**
 - 42 Active Policies
 - 18 AI Capabilities
 - 7 Modules
 - 4 Risk Levels
 - 5 Autonomy States
 - 6 Data Policy Groups
 - 3 Tenant Overrides
 - 2 Active Exceptions
- **Effective Governance Map**
- Visual map:

- **AI Capability**



- **Risk**



- **Autonomy**



- **Data Access**



- **Action Permission**



- **Approval Requirement**



- **Effective Decision**

Example

AI Configuration Assistant

Ticket Creation

Risk: Medium

Autonomy: L2

Result: Prepare

Pricing Change

Risk: High

Autonomy: L2

Result: Prepare + Approval

Access Security Change

Risk: Critical

Autonomy: L1

Result: Advisory Only

Recommendation AI

Product Ranking

Risk: Low

Autonomy: L4

Result: Auto-Allowed within guardrails

Modify Recommendation Strategy

Risk: Medium

Autonomy: L2

Result: Prepare

Change Price

Result:

- **Blocked**
- **Policy Lifecycle**

- **Draft**
- → **Review**
- → **Approval**
- → **Published**
- → **Effective**
- → **Superseded**
- → **Archived**
- **Version Comparison**
- Compare:
- **v1.3 → v1.4**
- Show:
 - Policies Added
 - Policies Removed
 - Autonomy Changes
 - Risk Changes
 - Data Access Changes
 - Approval Changes
 - Tenant Impact

Rollback

Allow rollback to a prior policy version where safe.

Governance should assess impact before rollback.

Actions

Compare Version | Submit for Approval | Publish | Schedule Publication | Assess Rollback

Critical Governance Architecture

Soft Labs should **not hard-code governance rules separately into each AI module**.

I recommend a centralized:

- **TICVAI AI Governance Policy Engine**

- Architecture:
- AI Capability



- AI Requested Operation



- Governance Context
 - └─ Tenant
 - └─ Venue
 - └─ User / Role
 - └─ Module
 - └─ Action
 - └─ Environment
 - └─ Data Categories

- Risk



- AI GOVERNANCE POLICY ENGINE



- | | |
|----------------------|--|
| Risk Classification | |
| Autonomy Policy | |
| Permission Policy | |
| Data Usage Policy | |
| Environment Policy | |
| Exception Policy | |
| Approval Requirement | |



Governance Decision



ALLOW

ALLOW WITH CONDITIONS

PREPARE ONLY

APPROVAL REQUIRED

ESCALATE

BLOCK

The requesting module then enforces the returned decision.

Governance Decision Object

Every governed operation should return something conceptually like:

- GovernanceDecision
- | |
|---------------|
| Decision ID |
| Capability ID |
| Tenant |
| User |
| Module |
| Action |
| Environment |
- | |
|--------------------|
| Risk Level |
| Autonomy Level |
| Effective Policies |
- | |
|----------|
| Decision |
|----------|

- | | — Allow
- | | — Allow with Conditions
- | | — Prepare Only
- | | — Approval Required
- | | — Escalate
- | — Block
- |
- | — Required Approvals
- | — Conditions
- | — Restrictions

| — Data Policy

| — Explanation

| — Policy Version

— Trace ID

This becomes extremely valuable later for Board 3 — AI Explainability & Audit.

Two-Layer Permission Model

A very important design principle for TICVAI:

AI permission should **never increase the user's permission**.

For example:

- User:
- Cannot modify pricing.
- AI:
- Has technical capability to call Pricing API.
- Result:

BLOCKED

The user cannot use AI as a way around RBAC/PBAC.

Therefore:

- **Effective AI Permission = User Permission $\cap$ AI Permission $\cap$ Governance Policy $\cap$ Module Policy**
- Whichever is most restrictive wins.

AI Autonomy Is Not the Same as User Permission

Soft Labs should also distinguish these concepts.

A user may be authorized to:

- **Modify Pricing**
- but AI autonomy may only be:
- **Level 2 — Prepare**
- Therefore:

User manually modifies price:

- Potentially permitted through normal Pricing workflow.

- User tells AI:
- “Change the price.”
- AI:

Prepares the proposed change and sends it for required approval.

This is intentional.

Policy Precedence

I recommend a clear precedence model.

Where policies conflict, the **more restrictive effective control should normally win**, subject to explicit governance design.

Conceptually:

- **Mandatory Platform Control**
- **Regulatory / Security Control**
- **Tenant Policy**
- **Venue / Business Restriction**
- **Module Policy**
- **Capability Policy**



Approved Exception

The exception mechanism must explicitly identify what it is allowed to override.

AI Provider Independence

Governance should apply regardless of whether TICVAI uses:

- • OpenAI models
- • Azure-hosted models
- • AWS AI services
- • Google AI
- • TICVAI-owned models
- • Customer-hosted models
- • Open-source models
- The governance question is not:
 - “Is this ChatGPT?”
 - It is:

What capability is being used, what data is involved, and what action can it perform?

Provider-specific technical controls will later belong in **AI Platform, Models & Providers**.

Policy Enforcement Must Be Server-Side

This is another critical requirement.

The frontend can show:

AI cannot execute this action.

But the real control must be enforced in backend services.

A UI button being disabled is **not governance**.

Every governed AI operation must be checked server-side before the owning module performs it.

Fail-Closed for Sensitive Actions

For High/Critical AI actions, if the Governance Policy Engine cannot determine permission because of:

- Policy service unavailable
- Policy conflict
- Missing risk classification
- Unknown capability
- Invalid user context
- TICVAI should generally **not execute the AI action**.
- Example:
 - Governance unavailable.
 - AI requests:
 - Modify production access rules.
 - Result:
 - **Execution Blocked**
 - This is safer than assuming permission.

Low-risk advisory capabilities may use different availability policies.

Important Boundary — Existing RBAC/PBAC

Do not rebuild TICVAI's existing authorization engine here.

RBAC/PBAC answers:

- Is this user permitted to perform this business action?
 - **AI Governance answers:**
 - Is AI permitted to participate in this action, at what autonomy level, and under what conditions?
 - They work together.
-

Important Boundary — AI Configuration Board 3

AI Configuration Board 3 already contains:

- Configuration approval
- Execution
- Rollback
- Change history

- We should **not duplicate those here**.
- AI Governance defines:

Pricing change = High Risk + Commercial Approval.

AI Configuration consumes that result and performs the actual configuration approval/execution workflow.

Important Boundary — Recommendation Board 5

Recommendation Board 5 may propose:

- Promote Recommendation Model v3.3.
- AI Governance determines:

Model promotion requires governance approval.

Recommendation Board 5 does not define the enterprise governance policy itself.

Important Boundary — AI Platform

Board 1 registers provider/model references for governance purposes, but it does not manage:

- • API Keys
- • Provider Credentials
- • Endpoint Configuration
- • Token Limits
- • Model Routing
- • Compute
- • Model Deployment

Those belong to the future **AI Platform, Models & Providers** module.

Board 1 — Final Screen Map

AI Governance & Policy Management

1. **AI Governance Command Center**
2. **AI Capability Registry & Ownership**
3. **AI Risk Classification & Assessment**
4. **AI Autonomy Level Configuration**
5. **AI Action & Permission Policy Builder**
6. **AI Data Access & Usage Policy**
7. **Environment, Tenant & Scope Governance**
8. **AI Policy Conflict, Exception & Override Management**
9. **AI Policy Testing & Governance Simulation**
10. **AI Governance Policy Publication & Effective Policy Map**

This gives Board 1 a very clear responsibility:

Define the rules of AI operation before we design how humans approve AI actions, how decisions are explained, or how AI governance risk is continuously monitored.

The next Board 2 should therefore be **AI Approval, Human Oversight & Exception Management**, focused on the human-in-the-loop operational workflow rather than redefining the policies created here.

AI Governance — Board 2: AI Approval, Human Oversight & Intervention Management

Module: AI Governance & Responsible AI Control
Board: 2 of 4
Screens: 10

Board Type: Human-in-the-Loop / Approval Orchestration / Oversight / Escalation / Intervention / Emergency Control

Board Purpose

Board 1 defined the **rules governing AI**:

- What may AI do, at what autonomy level, under which conditions, with which data, and at what risk level?
- Board 2 operationalizes those rules:

When human involvement is required, who needs to review the AI action, what information do they need, what decisions can they make, and how can humans intervene before or during AI execution?

The architecture should be:

- AI Request / Recommendation / Proposed Action



- Board 1 Governance Policy Engine



- Governance Decision



- If human oversight is required:
- Approval Requirement Resolution



- Approval Package



- **Human Review**



- **Approve / Conditional Approval / Request Change / Reject / Escalate**



- **Revalidation**



- **Owning Module Executes**



- **Oversight & Intervention**
- A major boundary is important:

Board 2 does not execute the underlying business operation.

For example, if AI proposes a price change, Board 2 manages the human approval. The **Pricing module** or AI Configuration execution orchestrator performs the actual approved change.

Screen 1 — AI Human Oversight Command Center

Purpose

Provide one central operational view of all AI activities requiring human oversight across TICVAI.

This should become the control room for AI actions waiting for human intervention.

Header KPIs

- Pending AI Approvals
- High-Risk Reviews
- Critical Reviews
- Awaiting Additional Approver
- Conditional Approvals
- Requested Changes
- Escalated Cases
- Overdue Reviews
- Active AI Executions Under Oversight
- Human Interventions Today

Oversight Queue

AI Capability	Proposed Action	Module	Risk	Status	Required By
Configuration AI	Modify Weekend Price	Pricing	High	Pending	Commercial Owner
Fraud AI	Restrict Account	Fraud	Critical	Human Review	Fraud Manager
Recommendation AI	Promote Model v3.3	Recommendation	Medium	Pending	AI Owner
Configuration AI	Publish New Product	Ticketing	Medium	Pending	Product Owner

Risk Distribution

Visual:

- **Low / Medium / High / Critical**
- **Approval SLA**
- Show:
 - • Within SLA
 - • Approaching SLA
 - • Overdue
 - • Escalated
- **AI Execution Oversight**
- Separate section:
- **Currently Executing Approved AI Actions**
- Example:
 - CFG-EXE-00218
 - Creating Dubai Museum configuration
 - Progress: **68%**
 - Status:

Human Oversight Available

Filters

- Tenant
- Venue
- AI Capability
- Module
- Risk
- Approver
- Status
- Environment
- Date
- SLA

Actions

Review Next | View Pending | View Executions | View Escalations | View History

Screen 2 — AI Approval Requirement & Routing Configuration

Purpose

Translate Board 1 governance decisions into the correct human approval workflow.

Board 1 may say:

- **Approval Required — Commercial Owner**
- Board 2 determines:

Who exactly is the Commercial Owner for this tenant/venue, whether additional approval is required, whether approval is sequential or parallel, and what happens if the approver is unavailable.

Approval Rule Inputs

Consume:

- • Governance Policy
- • Risk Level
- • AI Capability
- • Module
- • Action Type
- • Environment
- • Tenant
- • Venue
- • Financial Impact
- • Operational Impact
- • Security Impact
- • User Role
- **Approval Route Example**
- **AI Action**
- Modify Adult Weekend Price
- **Risk**
- High
- **Environment**
- Production
- **Route**
- **Step 1**
- Commercial Manager
- **Step 2**
- Venue General Manager



Approved for Execution

Approval Structures

Support:

- • Single Approval
- • Sequential Approval
- • Parallel Approval
- • Any-One-of-Group
- • All-of-Group
- • Conditional Approval
- • Risk-Based Approval
- • Amount / Impact-Based Approval
- **Example — Financial Threshold**
- Price adjustment:
- $\leq 5\%$
- → Commercial Manager

- **> 5% and ≤ 15%**
- → Commercial Manager + General Manager
- **> 15%**
- → Commercial Director + General Manager
- These are configurable examples, not hard-coded rules.
- **Fallback Approver**
- If primary approver is unavailable:
- **Delegate**
- or:
- **Escalate**
- according to policy.
- **Separation of Duties**
- Allow policies such as:

User who requested the AI change cannot approve the same change.

or:

AI Configuration Owner cannot be the sole approver for a Critical access-control modification.

Actions

Configure Route | Add Approver | Configure Threshold | Test Routing | Save

Screen 3 — AI Approval Review Workspace

Purpose

Give the human approver enough information to make an informed decision without having to navigate across many TICVAI modules.

This should be one of the strongest screens in the board.

Approval Header

Approval Request

AI-APR-000842

AI Capability

AI Configuration Assistant

Requested By

Venue Administrator

Environment

Production

Risk

HIGH

AI Proposed Action

Increase Adult Weekend Admission from AED 140 to AED 150 starting 1 October 2026.

Current vs Proposed

Configuration	Current	Proposed
Adult Weekend Price	AED 140	AED 150
Effective Date	Current	1 Oct 2026
B2C	AED 140	AED 150
POS	AED 140	AED 150

AI Reasoning / Recommendation

Show a controlled summary:

The requested change was interpreted as an effective-dated pricing modification affecting Adult Admission on weekend operating days.

Do not present hidden chain-of-thought.

Instead show:

- • User Request
- • AI Interpretation
- • Proposed Change
- • Relevant Decision Factors
- • Validation Result
- **Impact Analysis**
- Show:

Affected Products: 1

Affected Channels: 2

Existing Future Orders: 184

Existing Issued Tickets: Unchanged

Pricing Profiles: 1

Promotions Potentially Affected: 2

Governance Requirements

Risk: High

Policy: GOV-PRC-014

Required Approval:

Commercial Owner

Validation

- ✓ Pricing object valid
- ✓ Effective date valid
- ✓ Channel associations valid
- ✓ No blocking conflict
- ⚠ 2 promotions reference this product

Approver Actions

Approve

Approve with Conditions

Request Changes

Reject

Escalate

Screen 4 — Conditional Approval & Approval Conditions

Purpose

Allow humans to approve an AI action subject to explicit conditions rather than treating approval as simply Yes/No.

Example

AI proposes:

- Increase Weekend Price to AED 160.
- Approver decision:
 - **APPROVE WITH CONDITIONS**
 - Condition:
 - Maximum price AED 150.
 - The AI must not execute AED 160.
 - It must either:
 - • Modify the proposed action to AED 150 and revalidate, or
 - • Return to the requester for confirmation,
 - according to workflow policy.
- **Condition Types**
- Support:
 - • Value Limit
 - • Date Restriction
 - • Scope Restriction
 - • Venue Restriction
 - • Channel Restriction
 - • Product Restriction
 - • Time Restriction
 - • Volume Limit
 - • Additional Validation
 - • Additional Approval
 - • Manual Confirmation
 - • Custom Condition

Example

Original Proposal

Publish 12 products.

Conditional Approval

Approve 10 products. Keep VIP and Corporate products in Draft.

Machine-Readable Conditions

Where possible, approval conditions should be structured.

Example:

- MaximumPrice = AED 150
- EffectiveDate >= 01-Oct-2026
- Channels = B2C, POS

This allows the execution layer to verify compliance automatically.

Condition Validation

Before execution:

- **Approved Condition**

vs.

- **Execution Request**
- If different:
- **BLOCK EXECUTION**
- **Condition Expiry**
- Example:
- Approval valid until:
- **30 September 2026 23:59**
- After expiry:

Reapproval Required

Actions

Add Condition | Modify Condition | Approve with Conditions | Request Change

Screen 5 — Human Review, Challenge & AI Clarification Workspace

Purpose

Allow the approver to question the AI proposal before making a decision.

Human oversight should not mean simply clicking **Approve**.

Review Panel

Show:

- • Proposed Action
- • Business Context
- • Risk
- • Impact
- • AI Recommendation
- • Validation
- • Alternatives
- **Ask AI**
- Approver can ask:
- Why does this change affect the B2C and POS channels?
- AI should respond using the configuration evidence:

Both channels are associated with Pricing Profile PRC-3002, which contains the proposed Adult Weekend price.

Other Example Questions

Which existing promotions could be affected?

Will previously issued tickets change?

What happens if I reject this?

Can we apply this only to B2C?

Is there a lower-risk alternative?

AI Alternative

Example:

- **Original**
 - Change B2C + POS.
- **Alternative**
 - Change B2C only.
- **Impact:**
 - Lower scope.
- **Evidence References**
 - AI answers should reference:
 - Configuration Objects
 - Governance Policies
 - Validation Results
 - Impact Analysis
 - Existing Business Rules

Important Requirement

The AI must not fabricate explanations simply to satisfy the approver.

If evidence is unavailable:

- **Insufficient evidence available to determine this.**
- **Challenge Outcome**
 - The reviewer can:
 - Accept Explanation
 - Request Alternative
 - Modify Proposal
 - Return to Requester
 - Escalate

Actions

Ask AI | Compare Alternative | View Evidence | Request Modification | Continue Review

Screen 6 — Escalation, Delegation & Approval SLA Management

Purpose

Ensure AI approvals do not remain indefinitely pending and provide controlled routing when approvers are unavailable or additional authority is required.

Approval SLA Configuration

Example:

- **Medium Risk**
- 8 Business Hours
- **High Risk**
- 4 Business Hours
- **Critical**
- 1 Business Hour
- Configurable by tenant and policy.
- **Queue Status**

Request	Risk	Approver	Age	SLA	Status
APR-841	High	Commercial Manager	2h 12m	4h	Healthy
APR-842	High	GM	3h 48m	4h	Warning
APR-843	Critical	Security Owner	1h 14m	1h	Overdue

- **Escalation Flow**
- Example:
- **Commercial Manager**



- No response within 4 hours



- **Commercial Director**



- No response



- **General Manager**
- **Delegation**
- Support temporary authorized delegation.
- Example:
- **Primary Approver**
- Commercial Manager
- **Delegate**
- Deputy Commercial Manager
- **Valid**
- 15–22 September

Delegation Rules

Delegates must still possess the necessary RBAC/PBAC permission.

Delegation cannot bypass:

- Separation of duties
- Risk restrictions
- Mandatory role requirements
- **Escalation Reasons**
- SLA Expired
- Approver Unavailable

- Reviewer Requested Escalation
- Risk Increased
- Policy Conflict
- Critical Action
- Additional Expertise Required

Notifications

Board 2 should emit workflow notifications through TICVAI's notification capabilities rather than becoming a separate messaging engine.

Actions

[Escalate](#) | [Delegate](#) | [Reassign](#) | [Configure SLA](#) | [View Timeline](#)

Screen 7 — Live AI Execution Oversight & Human Intervention

Purpose

Allow authorized humans to monitor and intervene after an AI-driven action has been approved and entered execution.

This is different from approval.

Approval answers:

- May the action proceed?
- Oversight answers:
- What is happening while it proceeds, and can a human safely intervene?

- **Active Execution**

- Example:

- **Execution**

- AI-EXE-00218

- **Capability**

- AI Configuration Assistant

- **Plan**

- Dubai Museum Initial Setup

- **Environment**

- Production

- **Progress**

- 68%

- **Execution Steps**

Step	Action	Module	Status
1	Create Venue	Venue	Success
2	Create Products	Ticketing	Success
3	Create Schedule	Event	Success
4	Create Capacity	Capacity	Executing

Step	Action	Module	Status
5	Create Pricing	Pricing	Waiting
6	Create Media	Ticket Media	Waiting
7	Create Access	Access	Waiting

- **Human Controls**

- Depending on policy:
 - • Pause
 - • Resume
 - • Stop Remaining Actions
 - • Escalate
 - • Request Explanation
 - • Initiate Safe Rollback
 - • Take Manual Control

Example

Administrator notices:

- Capacity configuration is incorrect.
- Action:
- **PAUSE**
- Result:

Current safe operation finishes according to execution policy.

No downstream Pricing/Media/Access action starts.

Intervention Confirmation

High-impact intervention should show:

- **Impact of Pausing**
 - • Completed objects remain
 - • 4 actions will remain pending
 - • No active transaction affected
- **Emergency Stop**
- For appropriate Critical workflows:
- **Emergency Stop AI Execution**
- requires:
 - • Authorized Role
 - • Reason
 - • Confirmation
 - • Audit Record

Actions

Pause | Resume | Stop | Take Control | View Execution Trace

Screen 8 — Human Override & Manual Control Center

Purpose

Provide controlled mechanisms for humans to override an AI recommendation or take control when AI output is inappropriate.

Human override should always be possible where governance requires it.

Override Types

- Reject AI Recommendation
- Replace AI Recommendation
- Modify AI Proposal
- Stop AI Execution
- Manual Decision
- Manual Configuration
- Force Escalation
- Disable AI Capability Temporarily

Example — Recommendation

AI:

- Recommend Fast Pass.
- Authorized merchandising user:
- Override recommendation with Family Meal.
- The override should capture:
 - **Original AI Decision**
 - Fast Pass
 - **Human Decision**
 - Family Meal
 - **Reason**
 - Weekend family campaign
 - **Important**
 - The system must distinguish:
 - **Business Override**
 - from:

AI Failure

A legitimate merchandising override does not automatically mean the AI model was incorrect.

Example — Fraud

AI:

- High fraud risk — recommend account restriction.
- Human Fraud Analyst:
- Do not restrict.
- Reason:
- Verified legitimate corporate transaction.
- **Override Requirements**

- Depending on risk:
 - • Mandatory Reason
 - • Additional Approval
 - • Time Limit
 - • Scope
 - • Supporting Evidence
 - • Post-Review
- **AI Capability Pause**
- Authorized governance user may:
- **Pause Recommendation AI**
- Scope:
 - Dubai Theme Park
- Duration:
 - 2 hours
- Reason:
 - Unexpected recommendation behavior.
 - This should not necessarily disable AI platform-wide.

Actions

Override | Modify | Take Manual Control | Pause Capability | View Original AI Decision

Screen 9 — Approval & Intervention History / Decision Timeline

Purpose

Provide complete operational history of human involvement in AI activity.

This is the human-oversight timeline; Board 3 will later provide deeper enterprise AI explainability and audit.

Timeline Example

14:02

User requested AI pricing change.



14:02

AI prepared proposed configuration.



14:03

Governance classified High Risk.



14:03

Commercial approval requested.



14:18

Commercial Manager asked AI:

- Will existing tickets be affected?



- **14:19**

- AI provided evidence-based explanation.



- **14:22**

- Commercial Manager approved with condition:
- Maximum AED 150.



- **14:23**

- Configuration revalidated.



- **14:24**

- Execution authorized.



- **14:25**

- Pricing module applied effective-dated change.

- **History Table**

- Show:

- • Timestamp
- • Request
- • AI Capability
- • Human Actor
- • Role
- • Action
- • Previous Status
- • New Status
- • Reason
- • Approval Condition
- • Environment
- • Trace ID

Filters

- AI Capability
- User
- Approver
- Module
- Decision
- Risk
- Venue
- Tenant
- Date

Compare Human vs AI

For a selected case:

- **AI Proposal**
- AED 160
- **Human Approved**
- AED 150
- **Executed**
- AED 150
- This makes oversight effectiveness visible.

Actions

[View Timeline](#) | [View Proposal](#) | [View Approval](#) | [View Execution](#) | [Export History](#)

Screen 10 — Human Oversight Workflow Simulator & Readiness Center

Purpose

Allow Soft Labs/TICVAI administrators to test the entire human-in-the-loop workflow before activating it in production.

Scenario Builder

Example:

- **Capability**
- AI Configuration Assistant
- **User**
- Venue Administrator
- **Module**
- Pricing
- **Action**
- Increase Ticket Price
- **Change**
- AED 140 → AED 150
- **Environment**
- Production
- **Board 1 Governance Result**
- **Risk**
- High
- **Autonomy**
- Level 2 — Prepare
- **Governance Decision**
- Approval Required
- **Board 2 Routing**
- **Required Approver**
- Commercial Manager



- **Additional Approval**

- Not Required

Simulated Workflow

AI Prepares

✓



Validation

✓



Approval Request

✓



Commercial Manager

Pending



Approved



Revalidation

✓



Execution Authorization

✓



Owning Module

Pricing

Test Variations

Approver Rejects

Expected:

- **Execution Blocked**
- **Approver Requests Change**
- Expected:
- **Return to proposal/configuration workflow**
- **Approval Expires**
- Expected:
- **Reapproval Required**
- **Risk Changes High → Critical**
- Expected:
- **Approval Route Recalculated**
- **User Permission Removed**
- Expected:

- **Execution Blocked**
- **Governance Policy Changes While Pending**
- Expected:

Re-evaluate before execution

Readiness Checklist

- ✓ Approval route configured
- ✓ Approvers available
- ✓ Separation of duties valid
- ✓ SLA configured
- ✓ Escalation configured
- ✓ Conditional approval supported
- ✓ Intervention controls configured
- ✓ Audit enabled
- ✓ Revalidation configured

Overall Status

100% — Human Oversight Ready

Actions

Run Simulation | Test Failure | Test Escalation | View Effective Workflow | Activate

Critical Approval Architecture

Soft Labs should avoid building a completely separate approval engine specifically for AI.

TICVAI should use a shared enterprise approval/workflow capability, with AI Governance supplying the additional AI-specific context.

The architecture should be:

- AI Capability
- ↓
- Proposed AI Action
- ↓
- BOARD 1 — GOVERNANCE POLICY ENGINE
- ↓
- Governance Decision
- ↓
- Approval Required?
- ↓ YES
- Approval Requirement
- ↓
- BOARD 2 — HUMAN OVERSIGHT ORCHESTRATOR
- ↓
- Approval Route Resolution
- ↓
- Approval Package

- Human Review



- Approve / Conditional / Change / Reject / Escalate



- Revalidation



- Execution Authorization



OWNING TICVAI MODULE



Execution



Human Oversight / Intervention

AI Approval Package — Important Technical Object

Every human reviewer should receive a standardized package.

Conceptually:

- AIApprovalPackage
 - |
 - |— Approval ID
 - |— Governance Decision ID
 - |— AI Capability
 - |— Requested By
 - |— Tenant / Venue
 - |— Environment
 - |
 - |— Original User Intent
 - |— AI Interpretation
 - |— Proposed Action
 - |
 - |— Current State
 - |— Proposed State
 - |— Change Set
 - |
 - |— Risk Classification
 - |— Governance Policy
 - |— Required Approval
 - |
 - |— Impact Analysis

- └─ Validation Results
- └─ Dependencies
- └─ Warnings

|

└─ AI Recommendation / Explanation

└─ Evidence References

|

└─ Trace ID

The approver should never receive only:

“AI wants to modify pricing. Approve?”

They need enough information to understand the consequence.

Approval Is Bound to the Exact Change

This is a critical requirement.

Suppose the approver approves:

- Adult Weekend Price
AED 140 → AED 150
- AI must **not** subsequently execute:
- AED 140 → AED 160
- using the same approval.
- The approval should be bound to a:

Change Set / Configuration Version / Hash / Equivalent immutable reference.

If the proposal materially changes:

- **Approval becomes invalid**



- **Revalidation**



- **Reapproval**

- This prevents approval reuse.

Revalidation Before Execution

Approval alone is not sufficient.

Between approval and execution, something may change.

Example:

- 14:00
- Price change approved.
- 14:05
- Another administrator changes the Pricing Profile.
- 14:07
- AI attempts execution.

- TICVAI should detect:
 - **Underlying configuration changed since approval.**
 - Result:
 - **Execution Paused — Revalidation Required**
 - Depending on the difference:
 - • Continue
 - • Reapproval
 - • Rebuild Proposal
 - • Block
-

Risk Can Change During Workflow

Risk classification should not necessarily remain frozen forever.

Example:

- Initial AI request:
- Modify one product.
- Risk:
- **Medium**
- During preparation the system discovers:
- Change actually affects 420 products across 12 venues.
- Risk recalculated:

HIGH

The approval route should then be recalculated.

The previous lower-risk approval must not automatically authorize the expanded action.

Human-in-the-Loop vs Human-on-the-Loop

TICVAI should support both models.

Human-in-the-Loop

Human approval is required **before** AI execution.

Example:

Production pricing modification.

Human-on-the-Loop

AI may perform an approved low-risk category automatically, but authorized humans can monitor and intervene.

Example:

A Level 4 governed low-risk operation.

This is why **Screen 7 — Live AI Execution Oversight** is important even when individual approval is not required.

Human Override Must Not Destroy AI Learning

When a human overrides AI, record:

AI Decision

-

Human Decision

-

Reason

-

Final Outcome

Later, Recommendation/Fraud/AI analytics can determine whether human overrides tend to improve outcomes.

Do not simply overwrite the AI result in the database and lose the original decision.

Approval Delegation

Delegation must be permission-aware.

If:

- **Commercial Director**
- delegates approval to:

Operations Assistant

but Operations Assistant lacks Pricing approval rights,

the delegation should be:

- **INVALID**
 - Board 2 must still respect TICVAI RBAC/PBAC.
-

Approval Expiry

I strongly recommend supporting expiry.

Example:

- Pricing approval:
- **Valid for 24 hours**
- Configuration changed after that period:

Reapproval required

This avoids old approvals being used against a materially different operational situation.

AI Clarification vs Chain-of-Thought

For the UI, Soft Labs should not attempt to display hidden internal model reasoning.

Instead, TICVAI should provide structured explainability:

- **Request**
- → **Interpretation**
- → **Evidence**
- → **Decision Factors**

- → **Proposed Action**
 - → **Impact**
 - → **Governance Requirement**
 - This is both clearer and more defensible.
-

Important Boundary — Board 1

Board 1 defines:

- • Risk
- • Autonomy
- • AI permissions
- • Data usage
- • Policy
- • Governance decision
- Board 2 does not redefine those.

Board 2 consumes them to determine the required human workflow.

Important Boundary — Owning Modules

Board 2 does **not execute business actions**.

Examples:

- **Pricing change** → Pricing module
 - **Ticket publication** → Ticketing lifecycle
 - **Access restriction** → Access Control
 - **Recommendation model promotion** → AI Platform / Recommendation lifecycle
 - **Fraud restriction** → Fraud module
 - Board 2 provides:
 - **Execution Authorization**
 - not the business execution itself.
-

Important Boundary — Existing TICVAI Approval Workflow

Where TICVAI already has a central approval workflow/RBAC/PBAC framework, reuse it.

AI Governance should add AI-specific information:

- • AI Capability
 - • AI Risk
 - • AI Autonomy
 - • AI Proposal
 - • AI Explanation
 - • AI Evidence
 - • AI Trace
 - Do not create two unrelated approval infrastructures.
-

Important Boundary — Board 3

Board 2 records the human decisions needed for operational oversight.

The next board will go deeper into:

- **Board 3 — AI Explainability, Decision Trace & Audit**
- That board should answer:

What exactly happened from the original request through AI processing, governance, human decisions, model/provider interaction and final outcome—and can TICVAI prove it?

Board 2 — Final Screen Map

AI Approval, Human Oversight & Intervention Management

- 1. AI Human Oversight Command Center**
- 2. AI Approval Requirement & Routing Configuration**
- 3. AI Approval Review Workspace**
- 4. Conditional Approval & Approval Conditions**
- 5. Human Review, Challenge & AI Clarification Workspace**
- 6. Escalation, Delegation & Approval SLA Management**
- 7. Live AI Execution Oversight & Human Intervention**
- 8. Human Override & Manual Control Center**
- 9. Approval & Intervention History / Decision Timeline**
- 10. Human Oversight Workflow Simulator & Readiness Center**

This gives us a clean progression:

Board 1 — What is AI allowed to do?

- **Board 2 — Where must humans control or intervene?**
- **Board 3 — Can we explain and prove exactly what happened?**
- **Board 4 — Are AI risks, violations and governance health continuously under control?**

TICVAI — Scope of Work

AI Governance — Board 3: AI Explainability, Decision Trace & Audit



Board Purpose

Board 1 answered:

- What is AI allowed to do?
- Board 2 answered:
- When and how must humans approve, supervise, challenge, or intervene?
- Board 3 answers:

What exactly happened, why did TICVAI produce this AI outcome, what data/rules/models were involved, who approved or changed it, what was finally executed, and can we prove the complete history afterward?

This board gives TICVAI a **single end-to-end AI trace**.

The core architecture should be:

- **User / System Trigger**
- → **AI Context**
- → **Data & Evidence Used**
- → **Model / Rules / Tools Used**
- → **AI Output**
- → **Governance Evaluation**
- → **Human Decision**

- → **Execution**
- → **Business Outcome**
- → **Immutable Audit Trace**
- A key requirement:

TICVAI should provide **structured, evidence-based explainability**, not expose hidden model chain-of-thought.

Screen 1 — AI Explainability & Audit Command Center

Purpose

Provide management, governance teams and authorized auditors with one central view of AI decisions, explanations, trace completeness and audit status across TICVAI.

Header KPIs

- AI Decisions Today
- Governed AI Actions
- Automated Decisions
- Human-Reviewed Decisions
- Human Overrides
- Blocked Decisions
- Explainability Coverage
- Complete Decision Traces
- Incomplete Traces
- Audit Warnings

AI Activity Overview

Break down by:

- • Configuration AI
- • Recommendation AI
- • Forecasting AI
- • Fraud AI
- • AI Assistant
- • Dynamic Pricing AI
- • Future AI capabilities
- **Decision Status**
- Visual distribution:
- **Allowed**
- **Approval Required**
- **Human Modified**
- **Blocked**
- **Executed**
- **Failed**
- **Explainability Coverage**

- Example:
- **98.7% Explainability Coverage**
 - Complete Evidence: 97%
 - Partial Evidence: 2%
 - Missing Evidence: 1%
- **Recent AI Decisions**

Time	Capability	Decision / Action	Risk	Human	Result
14:32	Recommendation AI	Recommend Fast Pass	Low	No	Delivered
14:25	Configuration AI	Modify Weekend Price	High	Approved	Executed
14:18	Fraud AI	Restrict Account	Critical	Override	Not Restricted
14:05	Forecast AI	Attendance Forecast	Low	No	Generated

- **Trace Health**
- Show:
 - Complete
 - Partial
 - Missing Evidence
 - Provider Trace Missing
 - Execution Outcome Missing
 - Correlation Error

Filters

- Tenant
- Venue
- Capability
- Risk
- Model
- Provider
- Human Involvement
- Decision
- Environment
- Date

Actions

[Search Decision](#) | [Investigate Trace](#) | [View Exceptions](#) | [Export Audit](#) | [View Trace Health](#)

Screen 2 — AI Decision Explorer & Search

Purpose

Allow authorized users to locate any AI decision or AI-driven action across TICVAI.

This becomes the entry point when someone asks:

- “Why did TICVAI do this?”

- **Search By**

- Support:

- • Decision ID
- • Trace ID
- • Approval ID
- • Execution ID
- • Transaction ID
- • Customer Reference
- • Session ID
- • User
- • Product
- • Ticket
- • Order
- • Venue
- • AI Capability
- • Model Version
- • Date / Time

- **Example Search**

- **Transaction**

- ORD-928410

- Results:

Time	AI Capability	Decision	Related Object	Result
13:42	Recommendation AI	Fast Pass	ORD-928410	Accepted
13:44	Recommendation AI	Family Meal	ORD-928410	Ignored
13:46	Fraud AI	Risk Assessment	ORD-928410	Low Risk

- **Decision Card**

- Selecting a decision displays:

- **Decision ID**

- DEC-49102

- **Capability**

- Recommendation Engine

- **Customer**

- CUS-****821

- **Channel**

- B2C

- **Journey Stage**

- Checkout

- **Decision**

- Recommend Fast Pass

- **Rank**

- #1

- **Confidence**
 - 87%
- **Model**
 - REC-v3.2
- **Status**
 - Delivered
- **Relationship Search**
 - Allow navigation:
- **Customer**
 - → **Session**
 - → **Decision**
 - → **Delivery**
 - → **Transaction**

Sensitive Data

Display according to RBAC/PBAC and privacy policy.

An auditor may see:

- CUS-****821
- rather than full customer information.

Actions

[Open Decision](#) | [View Full Trace](#) | [View Related Decisions](#) | [Export Record](#)

Screen 3 — AI Decision Explanation Workspace

Purpose

Explain a selected AI decision in a clear business-readable format.

This should be usable by:

- • Administrator
- • Approver
- • Auditor
- • Operations Manager
- • Commercial Manager
- • AI Governance Team
- **Decision Summary**
 - Example:
 - **Why was Fast Pass recommended?**
 - **Decision**
 - Fast Pass
 - **Rank**
 - #1
 - **Confidence**
 - 87%
 - **Primary Decision Factors**

- Show structured factors such as:

- **Basket Context**

- Family Admission × 4

- Impact:

High Positive

Party Context

Family group detected from transaction context

Impact:

Positive

Product Relationship

Fast Pass configured as eligible upgrade/add-on

Impact:

Positive

Availability

Available

Impact:

Eligible

Capacity

Within configured operational threshold

Impact:

Eligible

Customer History

No recent Fast Pass purchase

Impact:

- Neutral

- **Hard Rules**

- Show separately:

✓ Product Active

✓ Channel Eligible

✓ Date Eligible

✓ Customer Eligible

✓ Capacity Available

✓ No Suppression

✓ Business Guardrails Passed

This distinction is important.

AI did not decide to ignore eligibility.

It ranked candidates **after deterministic controls passed**.

AI vs Rule Contribution

Example:

- **Eligibility**
- Rules
- **Candidate Generation**
- Hybrid
- **Ranking**
- AI
- **Diversity**
- Policy
- **Final Validation**
- Rules
- **Explanation Levels**
- Allow:

Simple

Fast Pass was recommended because it was eligible, available and highly relevant to the current family ticket basket.

Detailed

Show structured decision factors.

Technical

Show model/rule/feature/evidence references.

Important Requirement

Do not expose hidden reasoning tokens or chain-of-thought.

Actions

[View Evidence](#) | [View Candidate Comparison](#) | [View Technical Trace](#) | [Ask for Explanation](#)

Screen 4 — Data, Feature & Evidence Provenance

Purpose

Show exactly which data and evidence contributed to an AI decision and where each item originated.

This is critical for trust.

Evidence Table

Signal / Evidence	Value	Source	Time	Permission	Used
Basket Type	Family	Checkout	14:21	Allowed	✓
Party Size	4	Transaction	14:21	Allowed	✓
Membership	None	CRM	14:20	Allowed	✓

Signal / Evidence	Value	Source	Time	Permission	Used
Fast Pass Availability	Available	Capacity	14:21	Allowed	✓
Visit Date	Today	Ticketing	14:21	Allowed	✓
Payment Card	—	Payment	—	Blocked	✗

Signal Provenance

Selecting a signal shows:

- **Signal**
- Party Size = 4
- **Source System**
- Transaction Service
- **Source Object**
- Basket BKT-92018
- **Retrieved**
- 14:21:03.218
- **Freshness**
- Real Time
- **Permission Policy**
- DATA-AI-021
- **Derived Signals**
- Clearly label:
- **Source Data**

vs.

- **System Derived**

vs.

- **Model Inferred**
- Example:
- Likely Family
- Status:
- **Model Inferred**
- Confidence:
- 82%
- Do not display this as:

Confirmed Family Relationship

- unless TICVAI actually has confirmed relationship data.
- **Data Transformations**
- Where applicable:
- **Raw Value**



- **Masked / Normalized**



- **Feature**

- Example:
- Customer DOB



- Age Band



- AgeBand = Adult
- **Missing Data**
- Show:
- **Expected Signal**
- Membership Status
- **Status**
- Unavailable
- **Decision Handling**
- Neutral / fallback treatment

Never silently convert missing information into a negative preference.

Actions

[View Source](#) | [View Permission](#) | [View Transformation](#) | [Inspect Feature](#)

Screen 5 — Candidate, Rule & Decision Path Trace

Purpose

Explain how TICVAI moved from all possible options to the final AI decision.

This is particularly important for recommendation, fraud, pricing and configuration decisions.

Example — Recommendation

48 Potential Products



42 Active

6 removed



35 Channel Eligible

7 removed



28 Date / Time Eligible

7 removed



22 Customer Eligible

6 removed



18 Capacity Available

4 removed

**14 Passed Exclusions**

4 removed

**10 Passed Guardrails**

4 removed

**10 Ranked****Fast Pass #1****Candidate Table**

Candidate	Eligibility	AI Score	Rank	Final Result
Fast Pass	Eligible	0.91	1	Selected
Family Meal	Eligible	0.84	2	Selected
Annual Pass	Eligible	0.77	3	Placement Limit
VIP Tour	Ineligible	—	—	Capacity

Rule Trace

Show:

- **Rule**
 - CAP-014
 - Capacity must be available.
 - Result:
 - PASS
- **Exclusion Trace**
- **VIP Tour**
 - Excluded because:
 - Remaining Capacity = 0
 - Source:

Capacity Service

Configuration AI Example

For AI Configuration, the trace could instead show:

- **User Intent**



- **AI Interpretation**



- **Blueprint**



- **Configuration Plan**



Validation Rules



Governance



Approval



Execution

Fraud Example

Signals



Risk Rules



Model Score



Threshold



Recommended Action



Human Review

Actions

[View Rule](#) | [Compare Candidates](#) | [View Exclusion](#) | [View Decision Graph](#)

Screen 6 — Model, Provider & AI Runtime Trace

Purpose

Record which AI technology produced or contributed to the decision.

This gives TICVAI provider and model traceability without duplicating AI Platform configuration.

Runtime Information

Show:

- • AI Capability
- • Engine
- • Provider
- • Model
- • Model Version
- • Deployment Reference
- • Prompt / Template Version Reference
- • Rule Engine Version
- • Feature Set Version

- • Knowledge Source Version
- • Request Timestamp
- • Response Timestamp
- • Latency
- • Tokens / Compute where applicable
- • Provider Request Reference
- • Fallback Used
- • Retry Count

Example

Capability

AI Configuration Assistant

Provider

Approved Provider A

Model

Model X

Model Version

2026-09

Prompt Template

CFG-INTENT-v4.2

Knowledge Version

TICVAI-CONFIG-KG-v18

Latency

1.84 sec

Provider Routing

Show:

- **Primary Provider**



- Timeout



- **Fallback Provider**



- Success

This is important because the final result may not have been produced by the originally selected model.

Multi-Engine Decision

Example:

- Recommendation:
- **Rules Engine**
- Eligibility
-

ML Model

Ranking

-

Policy Engine

Guardrails

The trace should capture all contributors.

Version Traceability

If an auditor investigates six months later, TICVAI should know which model/version/rule set was active at the time.

Important Boundary

Configuration of providers, credentials and endpoints belongs to **AI Platform, Models & Providers**.

This screen only records and displays the runtime references.

Actions

[View Runtime](#) | [Compare Version](#) | [View Provider Trace](#) | [Open Model Reference](#)

Screen 7 — Governance, Approval & Human Decision Trace**Purpose**

Connect Board 1 and Board 2 governance activity to the AI decision.

Governance Timeline

Example:

- **14:02:12**
 - AI requested Pricing modification.
- ↓
- **14:02:13**
 - Governance Policy Engine evaluated request.
- ↓
- **Risk**
 - HIGH
- ↓
- **Autonomy**
 - Level 2 — Prepare
- ↓
- **Decision**
 - Approval Required
- ↓
- **14:03**
 - Commercial Manager notified.
- ↓

- **14:18**

- Approver requested clarification.



- **14:22**

- Approved with condition:
- Maximum AED 150.



- **14:23**

- AI proposal revalidated.



- **14:24**

- Execution authorized.

- **Governance Details**

- Show:

- • Governance Decision ID
- • Policy Version
- • Risk
- • Autonomy
- • Effective Policy
- • Approval Requirement
- • Exception if applicable
- • Human Reviewer
- • Human Decision
- • Conditions
- • Override
- • Delegation
- • Escalation

- **AI vs Human Decision**

- Display clearly:

- **Original AI Proposal**

- AED 160



- **Human Condition**

- Maximum AED 150



- **Final Authorized Change**

- AED 150

- **Human Override**

- Example:

- **AI Fraud Recommendation**

- Restrict Account

- **Human Decision**

- Do Not Restrict

- **Reason**
- Verified corporate transaction

Actions

[View Governance Policy](#) | [View Approval](#) | [View Human Decision](#) | [View Exception](#)

Screen 8 — Execution & Business Outcome Trace

Purpose

Connect the AI decision to what actually happened in TICVAI.

A good AI audit cannot stop at:

- “AI recommended X.”
- It needs to know:
 - **Was X actually executed, and what was the outcome?**
 - **Execution Trace**
 - Example:
 - **AI Proposal**
 - Change Adult Weekend Price



- **Approved**
- AED 150



- **Execution Authorization**
- AUTH-9281



- **Pricing API**
- Update Pricing Profile



- **Execution Result**
- SUCCESS
- **Before / After**
- **Before**
- Adult Weekend = AED 140
- **After**
- Adult Weekend = AED 150
- **Execution Metadata**
 - Execution ID
 - Owning Module
 - API / Service
 - Start Time
 - Completion Time
 - Result
 - Retry

- Failure
- Rollback
- Final State

Recommendation Outcome Example

Decision

Recommend Fast Pass



Delivered

Yes



Viewed

Yes



Accepted

Yes



Purchased

Yes



Revenue

AED 120

Important:

This is an observed outcome.

Board 5 Recommendation Analytics determines attribution and incrementality.

Failed Execution

Example:

- AI configuration approved.
- Pricing API failed.
- Trace:
- **Decision:** Approved
- **Execution:** Failed
- **Business State:** Unchanged
- This distinction must remain visible.
- **Rollback Trace**
- If rolled back:
- **Original Change**



- **Execution**



- Issue



- Rollback Authorization



- Rollback



- Restored State

Actions

[View Execution](#) | [View Before/After](#) | [View Outcome](#) | [View Rollback](#)

Screen 9 — AI Audit Record & Evidence Package

Purpose

Create a complete, tamper-evident audit package for an AI decision or group of decisions.

Audit Package Header

Audit Package

AI-AUD-2026-00918

Included Sections

1. Original Request / Trigger
2. User / System Context
3. AI Capability
4. Input Data References
5. Evidence Provenance
6. Rules Applied
7. Model / Provider / Version
8. AI Output
9. Governance Decision
10. Human Approval / Override
11. Execution Authorization
12. Execution Result
13. Business Outcome
14. Exceptions
15. Complete Timeline

Integrity Information

Record:

- • Audit Record ID
- • Created Timestamp
- • Source Trace IDs
- • Policy Version

- • Record Version
- • Integrity Verification
- • Retention Classification

Where the architecture supports cryptographic integrity, include hash/signature references.

Audit Scope

Allow:

- **Single Decision**
- **Customer Journey**
- **AI Capability**
- **Incident**
- **Date Range**
- **Model Version**
- **Governance Policy**

Export

Support controlled formats appropriate to TICVAI requirements, for example:

- • PDF
- • CSV
- • JSON / Machine-Readable Evidence Package
- Subject to permissions.
- **Sensitive Information**
- Export must respect:
 - • Data Access
 - • Privacy
 - • Masking
 - • Tenant Isolation
 - • Auditor Permission
- **Audit Package Status**
 - • Complete
 - • Partial
 - • Evidence Missing
 - • Under Investigation
 - • Archived

Actions

Generate Package | Verify Integrity | Export | Archive | Open Investigation

Screen 10 — AI Trace Investigation & Replay Simulator

Purpose

Allow authorized governance/technical teams to reconstruct an AI decision and understand whether the same conditions would produce the same or a different outcome.

This should be a powerful investigation tool.

Investigation Header**Decision**

DEC-49102

Original Time

13 Sep 2026 — 14:21:03

Original Decision**Context**

CTX-88210

Model

REC-v3.2

Policy

GOV-v1.4

Result

Fast Pass #1

Replay Modes**Exact Historical Replay**

Use preserved historical inputs/configuration where technically supported.

Current Configuration Replay

Use the same scenario against today's configuration.

Alternative Model Replay

Compare:

- Model v3.2

vs.

- Model v3.3
- **Alternative Policy Replay**
- Compare:
- Governance v1.4

vs.

- Governance v1.5
- **Comparison**

Element	Original	Replay
Model	v3.2	v3.3
Top Recommendation	Fast Pass	Family Meal
Confidence	87%	91%
Policy	v1.4	v1.4
Eligibility	Same	Same

- **Difference Analysis**
- Explain structured differences:

Candidate eligibility remained unchanged. Ranking changed because Model v3.3 assigned higher relevance to Family Meal under the preserved context.

Again, this should be evidence-based—not fabricated reasoning.

Investigation Notes

Authorized investigator can record:

- • Finding
- • Root Cause Category
- • Severity
- • Follow-Up Required
- • Related Incident
- • Assigned Owner

Important

A replay must **not execute production actions**.

It is simulation/investigation only.

Actions

Replay Decision | Compare Version | Compare Policy | Create Investigation | Close Review

Critical End-to-End Decision Trace

Soft Labs should implement one correlation model across all AI capabilities.

I recommend:

- AI Trace ID
- |
- |— Trigger / User Request
- |
- |— Context ID
- |
- |— Data & Evidence References
- |
- |— Rule / Feature References
- |
- |— AI Runtime
- | |— Provider
- | |— Model
- | |— Version
- |
- |— AI Decision ID
- |
- |— Governance Decision ID
- |
- |— Approval ID
- |

- |— Human Decision
- |
- |— Execution Authorization
- |
- |— Execution ID
- |
- |— Outcome References

The **Trace ID** should follow the process from beginning to end.

Standard AI Decision Record

Every significant TICVAI AI decision should conceptually create:

- AIDecisionRecord
- |
- |— Trace ID
- |— Decision ID
- |— Timestamp
- |
- |— Tenant
- |— Venue
- |— Environment
- |
- |— AI Capability
- |— Business Module
- |— Requested By / Triggered By
- |
- |— Intent / Task
- |— Context Reference
- |
- |— Evidence References
- |— Data Permissions
- |— Feature References
- |
- |— Rules Applied
- |— Candidates / Alternatives
- |
- |— Provider
- |— Model
- |— Model Version
- |— Runtime References
- |
- |— AI Output



Not every AI chat response needs the same heavy audit depth. The level of trace should be configurable based on risk and capability.

Explainability Levels

I recommend three explainability levels.

Level 1 — Business Explanation

For normal operational users.

Example:

Fast Pass was recommended because it is available, eligible for this basket and highly relevant to the current visit context.

Level 2 — Governance Explanation

For approvers/managers.

Includes:

- • Key Factors
- • Eligibility
- • Rules
- • Risk
- • Confidence
- • Policy
- • Alternatives

- **Level 3 — Technical / Audit Trace**

- For authorized governance and technical teams.
- Includes:

- • Model Version
 - • Feature References
 - • Data Provenance
 - • Rules
 - • Provider
 - • Latency
 - • Correlation IDs
 - • Execution References
 - • Complete timeline
 - RBAC/PBAC determines access.
-

Explainability Must Match the AI Type

Soft Labs should not design one generic explanation for every AI capability.

Recommendation AI

Explain:

- **Why this candidate ranked above others.**
- **Forecasting AI**
- Explain:

Which major variables/signals contributed to the forecast and what uncertainty exists.

Fraud AI

Explain:

Which permitted risk signals contributed to the risk result and what rules/thresholds applied.

Configuration AI

Explain:

How the user request was interpreted and translated into the proposed configuration.

AI Assistant

Explain:

Which approved knowledge/data sources supported the response where applicable.

Confidence Is Not the Same as Accuracy

This distinction should be reflected in the UI.

If AI says:

- **Confidence 92%**
- that does not mean:

“The answer is 92% guaranteed correct.”

Confidence must be interpreted according to the particular model/capability and calibration method.

Board 3 should therefore display:

- **Confidence / Model Score**
- rather than misleading:
- **Accuracy = 92%**

- unless it genuinely represents measured accuracy.

Audit Immutability

Once a governed AI action has occurred, the historical record should not be editable like a normal configuration object.

If correction is necessary:

- **Original Record**
 - remains
 -

Correction / Annotation Record

is added.

Do not overwrite history.

Retention

Retention should be configurable based on:

- • Tenant Policy
- • Data Category
- • Risk
- • Capability
- • Regulatory/Contractual Requirement
- • Environment

For example, Critical governance decisions may require different retention than low-risk anonymous recommendation traces.

Board 3 controls AI audit classification and evidence linkage; broader enterprise retention/privacy policy remains authoritative.

Tenant Isolation

Because TICVAI is multi-tenant, an AI auditor from:

- **Tenant A**
 - must never search or inspect:

Tenant B AI traces

unless there is an explicitly authorized TICVAI platform-level role.

This must be enforced server-side.

Audit Integrity

For important governed AI records, TICVAI should support tamper-evident integrity.

Conceptually:

Decision Record

-

Timestamp

-

Evidence References

-

Governance Decision

-

Execution Result**Integrity Record**

This can later support stronger technical mechanisms without making Board 3 dependent on a particular implementation.

Replay Limitation

One important point for Soft Labs:

An LLM or stochastic model may not reproduce exactly the same text/output when rerun.

Therefore **Exact Historical Replay** should primarily rely on the **preserved original request, inputs, output, model/version and runtime trace**.

A new model call is a **re-simulation**, not proof of what originally happened.

The UI should distinguish:

- **Historical Reconstruction**
- from:

New Simulation**Important Boundary — Board 1**

Board 1 owns:

- Governance Policies
- Risk
- Autonomy
- **Permissions**
- Data Usage Rules

Board 3 only displays the exact policy/version that applied to the historical decision.

Important Boundary — Board 2

Board 2 owns:

- Human Approval
- Human Challenge
- Intervention
- Override
- Escalation

Board 3 consumes those records to build the complete trace.

Important Boundary — Recommendation Board 5

Board 3 can show:

- Fast Pass recommendation → purchased → AED 120.
- But it should **not claim incremental revenue**.
- Recommendation Board 5 owns:
 - • Attribution
 - • Incrementality
 - • Experiments
 - • Recommendation Performance

Important Boundary — AI Platform

Board 3 displays:

- • Provider
- • Model
- • Version
- • Runtime Reference
- • Latency
- • Fallback
- But AI Platform owns actual:
 - • Provider Setup
 - • Credentials
 - • Model Deployment
 - • Routing
 - • Endpoint Management
 - • Cost Controls

Important Boundary — Enterprise Audit

TICVAI's general enterprise audit may capture:

- User X changed Product Y.
- Board 3 adds the specialized AI dimension:

User requested AI → AI interpreted request → Model X contributed → Governance Policy Y applied → Human Z approved → Service executed → Business state changed.

These should integrate rather than create disconnected audit systems.

Board 3 — Final Screen Map

AI Explainability, Decision Trace & Audit

1. **AI Explainability & Audit Command Center**
2. **AI Decision Explorer & Search**
3. **AI Decision Explanation Workspace**

4. **Data, Feature & Evidence Provenance**
5. **Candidate, Rule & Decision Path Trace**
6. **Model, Provider & AI Runtime Trace**
7. **Governance, Approval & Human Decision Trace**
8. **Execution & Business Outcome Trace**
9. **AI Audit Record & Evidence Package**

10. AI Trace Investigation & Replay Simulator

This makes the first three governance boards very clean:

- **Board 1 — CONTROL**

What is AI allowed to do?



- **Board 2 — HUMAN OVERSIGHT**

Where and how must people approve or intervene?



Board 3 — EXPLAIN & PROVE

Why did the AI do it, what exactly happened, and can TICVAI reconstruct the complete evidence trail?

The final **Board 4** should then become the ongoing operational layer: **AI Risk, Compliance & Governance Monitoring**—policy violations, AI incidents, drift/governance alerts, compliance health, control effectiveness and executive governance reporting.

AI Governance — Board 4: AI Risk, Compliance & Governance Monitoring

Module: AI Governance & Responsible AI Control

Board: 4 of 4

Screens: 10

Board Type: Continuous Governance Monitoring / Risk / Compliance / Control Effectiveness / Incidents / Alerts / Executive Oversight



Board Purpose

Board 4 is the **continuous monitoring and assurance layer** of TICVAI AI Governance.

The previous boards established:

- **Board 1 — CONTROL**
What is AI allowed to do?
- **Board 2 — HUMAN OVERSIGHT**
Where must humans approve, supervise or intervene?

Board 3 — EXPLAIN & PROVE

Why did AI produce the outcome, and can TICVAI reconstruct the complete evidence trail?

Board 4 answers:

Is AI continuing to operate safely, compliantly and according to the governance rules we approved—and what should TICVAI do when something goes wrong?

The continuous loop should be:

- **AI Activity**
- → **Governance Telemetry**
- → **Policy Compliance**
- → **Risk Signals**
- → **Control Effectiveness**

- → **Anomaly / Violation Detection**
- → **Alert**
- → **Investigation**
- → **Incident / Remediation**
- → **Governance Review**
- → **Policy / Control Improvement**

Screen 1 — AI Governance Monitoring Command Center

Purpose

Provide executives, governance teams and authorized administrators with one centralized real-time view of AI governance health across TICVAI.

Header KPIs

- Active AI Capabilities
- AI Decisions Today
- Governance Compliance Rate
- Policy Violations
- High-Risk Events
- Critical AI Events
- Open AI Incidents
- Active Exceptions
- Controls Passing
- Controls Failing
- AI Capabilities Under Review
- Governance Health Score

Governance Health

Example:

- **AI Governance Health — 94%**
- Breakdown:
 - Policy Compliance — 98%
 - Approval Compliance — 99%
 - Data Usage Compliance — 97%
 - Trace Completeness — 98%
 - Control Effectiveness — 91%
 - Open Risk Exposure — 88%

The score methodology must be transparent and configurable rather than an unexplained AI-generated number.

Capability Health

Capability	Risk	Compliance	Alerts	Incidents	Status
Configuration AI	High	99.2%	2	0	Healthy

Capability	Risk	Compliance	Alerts	Incidents	Status
Recommendation AI	Medium	98.1%	4	0	Monitor
Dynamic Pricing AI	High	100%	0	0	Healthy
Fraud AI	Critical	96.4%	3	1	Review
Forecasting AI	Low	99.8%	0	0	Healthy

Risk Heatmap

Dimensions:

- **AI Capability × Risk Category**

- Categories:

- • Policy
- • Data
- • Security
- • Model
- • Human Oversight
- • Operational
- • Customer
- • Financial
- • Compliance

- **Active Critical Items**

- Example:

CRITICAL

Fraud AI — unexpected increase in account restriction recommendations.

HIGH

Recommendation AI — fallback rate exceeded threshold.

MEDIUM

Configuration AI — approval SLA degradation.

Filters

- Tenant
- Venue
- Capability
- Module
- Environment
- Risk
- Status
- Provider
- Model
- Date

Actions

[Investigate Risk](#) | [View Alerts](#) | [View Incidents](#) | [View Controls](#) | [Governance Review](#)

Screen 2 — AI Risk Register & Risk Exposure Management

Purpose

Maintain the enterprise risk register specifically for TICVAI AI capabilities.

Board 1 classifies individual capabilities/actions.

Board 4 manages the **ongoing risk exposure** after those capabilities are operational.

Risk Register

Each risk record contains:

- • Risk ID
- • Risk Name
- • AI Capability
- • Module
- • Risk Category
- • Description
- • Inherent Risk
- • Existing Controls
- • Residual Risk
- • Risk Owner
- • Treatment
- • Review Date
- • Status

Example

Risk

AI Recommendation Fatigue

Capability

Recommendation Engine

Category

Customer Experience

Inherent Risk

High

Existing Controls

- Frequency Caps
- Suppression Rules
- Journey Limits
- Customer Rejection Handling

Residual Risk

Medium

Risk Categories

Support:

- • Customer
- • Commercial
- • Financial
- • Operational
- • Security
- • Privacy / Data
- • Compliance
- • Model / AI Quality
- • Explainability
- • Human Oversight
- • Reliability
- • Third-Party / Provider
- • Reputation
- **Risk Matrix**
- Show:
- **Likelihood × Impact**
- with:
 - • Low
 - • Medium
 - • High
 - • Critical
- **Risk Treatment**
- Options:
- **Accept**
- **Mitigate**
- **Transfer**
- **Avoid**
- **Monitor**
- **Risk Trend**
- Show:
- **Improving**
- **Stable**
- **Deteriorating**
- **Risk Linkage**
- Risk can link to:
 - • AI Capability
 - • Governance Policy
 - • Control
 - • Alert
 - • Incident
 - • Exception

- Model
- Provider

Actions

Create Risk | Assess Risk | Add Control | Assign Owner | Review Risk

Screen 3 — AI Governance Control Library & Control Effectiveness

Purpose

Define and continuously test whether the controls designed to govern AI are actually working.

A policy existing on paper does not mean the control is effective.

Control Library

Examples:

- **AI-GOV-C001**
 - Production Pricing AI Requires Approval
- **AI-GOV-C002**
 - Critical AI Actions Cannot Auto-Execute
- **AI-GOV-C003**
 - AI Cannot Exceed User RBAC/PBAC
- **AI-GOV-C004**
 - Customer Payment Data Cannot Be Sent to Recommendation AI
- **AI-GOV-C005**
 - Every High-Risk AI Action Requires Complete Trace
- **Control Record**
 - Control ID
 - Control Name
 - Objective
 - Risk Addressed
 - Policy Reference
 - AI Capabilities
 - Control Type
 - Frequency
 - Test Method
 - Owner
 - Last Test
 - Result
 - Next Test
- **Control Types**
 - Preventive
 - Detective
 - Corrective
 - Manual

- Automated
- Hybrid

Example Test

Control

Critical AI actions require human approval.

Test Population

2,418 Critical AI actions

Result

2,418 required approval.

Effectiveness: 100%

Failed Control Example

Control

High-risk AI traces must include provider/model version.

Population:

- 4,210
- Complete:
- 4,183
- Missing:
- 27
- Effectiveness:
- **99.36%**
- Status:
- **Review Required**
- **Control Health**
- Statuses:
 - Effective
 - Partially Effective
 - Ineffective
 - Not Tested
 - Overdue
 - Disabled

Actions

[Create Control](#) | [Run Test](#) | [View Evidence](#) | [Create Remediation](#) | [Schedule Test](#)

Screen 4 — AI Policy Compliance & Violation Monitoring

Purpose

Continuously detect AI activity that violates or attempts to violate Board 1 governance policies.

Compliance KPIs

- Evaluated Actions

- Compliant
- Blocked by Policy
- Policy Violations
- Attempted Violations
- Exception-Based Actions
- Unknown / Unclassified Actions

Violation Types

Examples:

- • Autonomy Exceeded
- • Unauthorized AI Action
- • User Permission Violation
- • Data Usage Violation
- • Approval Bypass Attempt
- • Environment Restriction Violation
- • Expired Exception Used
- • Unregistered AI Capability
- • Missing Governance Evaluation
- • Policy Version Mismatch
- **Example Violation**
- **Capability**
- AI Configuration Assistant
- **Requested Action**
- Modify Access Security Policy
- **Environment**
- Production
- **Permitted Autonomy**
- Level 1 — Advisory
- **Attempted**
- Execute
- **Result**
- **BLOCKED**
- **Violation Severity**
- **Critical**

Policy

GOV-ACCESS-021

Actor

Venue Administrator + AI Capability

Violation Timeline

Show:

- **Request**
- → **Governance Evaluation**
- → **Violation Detected**

- → **Execution Blocked**
- → **Alert Created**
- **Repeated Violation Detection**
- Example:
- Same capability attempted prohibited action:
- **14 times in 24 hours**
- Status:

Escalated

Actions

[Investigate](#) | [View Trace](#) | [Create Incident](#) | [Suppress False Positive](#) | [Review Policy](#)

Screen 5 — AI Data, Privacy & Usage Compliance Monitoring

Purpose

Monitor whether AI capabilities are actually using data according to the policies configured in Board 1.

This screen is monitoring—not the configuration of the data policies themselves.

Data Usage KPIs

- AI Data Requests
- Allowed
- Conditionally Allowed
- Blocked
- External Transfers
- Masked / Redacted
- Data Policy Violations
- Unknown Data Categories

Data Usage Matrix

Capability	Data Category	Usage	External	Policy	Status
Recommendation AI	Basket	Process	Yes*	DATA-021	Compliant
Recommendation AI	Payment Card	Blocked	No	DATA-024	Compliant
Forecast AI	Attendance	Process	Yes*	DATA-031	Compliant
Config AI	Customer Profile	Attempted	No	DATA-018	Blocked

* only where provider/data policy allows it.

External Provider Monitoring

Show:

- **Provider**
- Approved Provider A
- **Data Categories Sent**
 - Product

- • Basket Context
- • Venue
- **Prohibited Data**
 - • Full Payment Details
 - • Restricted PII
- **Data Minimization Monitoring**
- Compare:
- **Requested Data**

vs.

- **Actually Sent**
- Example:
- 24 available fields



- 8 required fields



- 6 sent after masking
- **Sensitive Data Alert**
- Example:

HIGH

Unexpected customer identifier detected in an AI provider payload.

Response:

- **Payload Blocked**
- **Consent / Permission Status**
- Where customer data is used:
 - Consent / lawful-use reference where applicable
- **Purpose**
- Data Policy
- Retention
- Provider Restriction

Actions

[Investigate Usage](#) | [View Payload Metadata](#) | [View Data Policy](#) | [Create Incident](#) | [View Provider](#)

Screen 6 — AI Quality, Behavior & Governance Drift Monitoring

Purpose

Detect meaningful changes in AI behavior that may increase governance risk.

This is not the full model monitoring platform.

Board 4 focuses specifically on **governance-relevant behavioral changes**.

Monitoring Signals

Examples:

- • Decision Distribution Shift
- • Confidence Distribution Shift
- • Human Override Rate
- • Rejection Rate
- • Fallback Rate
- • Approval Rate
- • Block Rate
- • Recommendation Diversity
- • Fraud Flag Rate
- • Configuration Error Rate
- • Explanation Availability
- • Decision Latency

Example

Recommendation Engine v3.3

Normal:

- Fast Pass recommended to 18–24% of eligible customers.
- Current:
- Fast Pass = **41%**
- Change:
- +86%
- Status:
- **Behavior Drift Warning**
- **Another Example**
- **Fraud AI**
- Historical Restriction Recommendation Rate:
- 1.2%
- Current:
- 4.8%
- Status:
- **HIGH RISK CHANGE**
- **Override Drift**
- Human Override Rate:
- Historical:
- 4%
- Current:
- 17%
- This may indicate:
 - • Model quality issue
 - • Changed business conditions
 - • Policy mismatch
 - • New customer behavior
 - • Incorrect thresholds
- TICVAI should not automatically assume the cause.

- **Threshold Configuration**
- Monitor using:
 - Absolute Threshold
 - Percentage Change
 - Standard Deviation / Statistical Method
 - Model-Specific Metric
 - Business Baseline
 - Custom Rule
- **Model / Version Comparison**
- Show behavior before and after:

v3.2 → v3.3

Actions

[Investigate Drift](#) | [Compare Version](#) | [View Decisions](#) | [Create Alert](#) | [Open Model Review](#)

Screen 7 — AI Governance Alert & Detection Center

Purpose

Centralize governance-related alerts generated from policy, control, data, behavior and operational monitoring.

Alert Sources

- Policy Violation
- Control Failure
- Data Violation
- Behavioral Drift
- Approval Issue
- Audit Trace Failure
- Provider Issue
- Security Event
- Exception Expiry
- Unregistered AI Activity
- Excessive Human Override
- AI Service Reliability

Alert Record

- Alert ID
- Timestamp
- Capability
- Tenant / Venue
- Source
- Severity
- Description

- Evidence
- Related Policy
- Related Risk
- Status
- Owner

Severity

- Informational
- Low
- Medium
- High
- Critical

Alert Correlation

Very important:

Do not create 500 separate alerts for the same underlying problem.

Example:

- 427 Recommendation AI fallback events



- Correlated into:
- **One Governance Alert**
- **Recommendation AI Fallback Rate Exceeded**
- **Alert Lifecycle**
- **Detected**
- → **Triaged**
- → **Assigned**
- → **Investigating**
- → **Resolved**
- → **Closed**
- or:
- → **Converted to Incident**
- **False Positive**
- Authorized reviewer may classify:
- **False Positive**
- with reason.

This should tune monitoring configuration where appropriate but not erase the historical alert.

Actions

Acknowledge | Assign | Investigate | Create Incident | Close

Screen 8 — AI Incident & Remediation Management

Purpose

Manage significant AI governance failures from detection through containment, investigation, remediation and closure.

Incident Example

Incident

AI-INC-00281

Capability

Fraud Intelligence

Severity

Critical

Issue

Unexpected increase in customer account restriction recommendations after model update.

Incident Timeline

10:02

Model v4.2 activated.



10:18

Restriction rate exceeded threshold.



10:19

Governance alert generated.



10:23

Fraud Manager acknowledged.



10:27

AI capability restricted to Advisory Mode.



10:32

Investigation started.

Containment Actions

Depending on authorization:

- Reduce Autonomy
- Pause Capability
- Disable Model Version
- Switch to Fallback
- Require Human Review

- Restrict Data Access
- Disable Specific Action
- Revert Policy / Model where authorized
- **Root Cause Categories**
 - Model
 - Data
- **Configuration**
 - Policy
 - Provider
 - Integration
 - Human Error
 - Business Change
 - Unknown

Remediation Plan

Action	Owner	Due	Status
Review v4.2 behavior	AI Team	14 Sep	In Progress
Restore v4.1	Platform	Immediate	Complete
Revalidate thresholds	Fraud Team	15 Sep	Pending
Run control test	Governance	15 Sep	Pending

Post-Incident Review

Capture:

- Root Cause
- Impact
- Customers / Transactions Affected
- Control Failure
- Corrective Actions
- Preventive Actions
- Lessons Learned

Closure Requirements

Critical incident cannot close until required remediation is verified.

Actions

[Contain](#) | [Assign](#) | [Add Remediation](#) | [Link Audit Trace](#) | [Close Incident](#)

Screen 9 — AI Compliance, Assurance & Governance Reporting

Purpose

Provide structured governance evidence and management reporting without duplicating TICVAI's generic BI platform.

Governance Reports

Support predefined reports such as:

- • AI Capability Inventory
- • AI Risk Register
- • AI Policy Compliance
- • Human Oversight Compliance
- • AI Approval Compliance
- • AI Data Usage
- • AI Exceptions
- • AI Incidents
- • Control Effectiveness
- • AI Explainability Coverage
- • Model / Provider Governance
- • High-Risk AI Activity

- **Governance Assurance Dashboard**

- Example:

Governance Area	Target	Actual	Status
High-Risk Approval	100%	100%	Pass
Critical Auto-Execution	0	0	Pass
Trace Completeness	99%	98.7%	Review
Data Policy Compliance	100%	99.9%	Review
Control Testing On Time	100%	96%	Warning

- **Reporting Period**

- • Daily
- • Weekly
- • Monthly
- • Quarterly
- • Custom

- **Tenant Governance Report**

- Allow per-tenant reporting.

- Example:

- **Dubai Theme Park**

- AI Governance Report
- September 2026

- **Evidence Links**

- Every KPI should be drillable to:

- **Metric**

- → **Underlying Records**

- → **Trace**

- → **Evidence**

- This is much stronger than a static dashboard.

- **Report Generation**
- Support controlled exports where required.

Actions

[Generate Report](#) | [View Evidence](#) | [Schedule Report](#) | [Export](#) | [Submit for Review](#)

Screen 10 — AI Governance Review, Action Plan & Continuous Improvement

Purpose

Bring all governance monitoring together into a structured periodic review and improvement cycle.

This is the final screen of the entire AI Governance module.

Governance Review

Example:

- **September 2026 AI Governance Review**
- **Overall Status**
- **Healthy with 3 Areas Requiring Action**
- **Summary**
- **Policy Compliance**
- 98.9%
- **Open High Risks**
- 3
- **Critical Incidents**
- 1
- **Control Effectiveness**
- 96.8%
- **Human Oversight Compliance**
- 99.7%
- **Trace Completeness**
- 98.7%
- **Priority Findings**
- **1. Recommendation AI Fallback Rate**
- Status:
- Elevated
- Recommendation:

Review model/provider reliability.

2. Fraud AI Override Rate

Status:

- Above baseline
- Recommendation:

Review decision thresholds and recent model version.

3. AI Trace Completeness

Status:

- Below 99% target
- Recommendation:
- Resolve missing provider version references.
- **Governance Improvement Plan**

Finding	Action	Owner	Priority	Due
Fallback Rate	Provider Review	AI Platform	High	18 Sep
Fraud Override	Model Review	Fraud Team	High	16 Sep
Missing Trace	Integration Fix	Platform	Medium	20 Sep

- **AI-Generated Governance Recommendations**

- AI may assist by suggesting:

“Human override rate increased following Model v4.2 deployment. Consider reviewing v4.2 before expanding traffic.”

But this must be presented as:

AI Recommendation

not an established root cause unless evidence proves causation.

Improvement Actions

May create controlled requests to:

- **Board 1**
- Change governance policy.
- **Board 2**
- Strengthen approval/oversight.
- **Board 3**
- Improve trace/evidence requirements.
- **AI Platform**
- Review provider/model.
- **Owning Module**
- Modify operational configuration.
- **Review Approval**
- Governance review can be:
- **Draft**
- → **Reviewed**
- → **Approved**
- → **Actions Assigned**
- → **Completed**

Actions

Create Action Plan | Assign Owners | Submit Review | Track Actions | Close Review

Critical Continuous Monitoring Architecture

I recommend Soft Labs implement Board 4 around a central **AI Governance Telemetry & Monitoring Layer**.

AI CAPABILITIES

- |
- |— Configuration AI
- |— Recommendation AI
- |— Forecasting AI
- |— Fraud AI
- |— AI Assistant
- |— Future AI



AI GOVERNANCE TELEMETRY



Governance Decisions	
Policy Evaluations	
AI Decision Traces	
Human Approvals	
Overrides	
Data Usage	
Model / Provider Runtime	
Execution Outcomes	
Exceptions	



GOVERNANCE MONITORING ENGINE



Policy Compliance	
Risk Monitoring	
Control Testing	
Data Compliance	
Behavior / Drift Signals	
Alert Detection	



Alert



Investigation



Incident / Remediation



Governance Review



Controlled Improvement

Governance Telemetry Event

Soft Labs should standardize monitoring events.

Conceptually:

- GovernanceTelemetryEvent
 - |
 - |— Event ID
 - |— Trace ID
 - |— Timestamp
 - |
 - |— Tenant
 - |— Venue
 - |— Environment
 - |
 - |— AI Capability
 - |— Module
 - |— Model / Provider Reference
 - |
 - |— Event Type
 - |— Risk Level
 - |— Governance Policy
 - |
 - |— Governance Decision
 - |— Human Decision
 - |
 - |— Data Categories
 - |— Execution Result
 - |
 - |— Control References

|— Exception Reference

|

|— Monitoring Metadata

This allows Board 4 to monitor all AI capabilities consistently.

Policy Violation vs Blocked Attempt

This distinction is important.

Suppose AI attempts something prohibited.

Board 1 correctly blocks it.

That is:

- **Blocked Attempt**
 - The governance system worked.
 - But if AI actually performs the prohibited action:
- **Policy Violation**
 - That is much more serious.
 - The dashboards should not treat both as identical.
 - Example:
 - **14 prohibited attempts**
 - **14 blocked**
 - Governance control:
- **Effective**
 - versus:
 - **14 prohibited attempts**
 - **13 blocked**
 - **1 executed**
 - Governance control:

FAILED

Control Effectiveness Is Critical

This board should measure not just:

- “Do we have a policy?”
 - but:
 - “Did the policy actually prevent or detect the risk?”
 - Therefore:
 - **Policy**
 - → defines requirement.
 - **Control**
 - → implements requirement.
 - **Control Test**
 - → proves whether it worked.
 - **Evidence**
 - → supports the test.
-

AI Risk Should Be Dynamic

Board 1 provides the initial risk classification.

Board 4 should monitor whether operational evidence suggests the risk needs review.

Example:

- Recommendation AI originally:
- **Medium Risk**
- But later:
 - • Very high customer reach
 - • Increased human overrides
 - • Multiple incidents
 - • New sensitive data usage
- TICVAI may flag:

Risk Classification Review Recommended

It should not silently change the official risk level unless governance policy allows it.

Incident Containment

For serious AI incidents, Board 4 should be able to request rapid governance actions.

Example:

- **Fraud AI issue detected**



- Set capability:

Level 4 → Level 1 Advisory

This should go through a controlled emergency governance mechanism.

Another option:

- **Pause Specific Model**
- without disabling the entire AI platform.

This requires integration with the future AI Platform module.

Monitoring Must Not Depend on the Same AI Being Monitored

An important architectural point:

Core governance controls should remain deterministic wherever possible.

For example:

“Did Critical AI action receive mandatory approval?”

should be evaluated from structured system records—not by asking an LLM.

AI can help investigate and summarize governance information, but it should not be the sole judge of whether its own mandatory control was followed.

AI-Generated Governance Recommendations

AI can still add significant value.

For example:

- “Recommendation AI human override rate increased from 4% to 13% after deployment of model v3.4.”

- Then:
- “Consider comparing v3.3 and v3.4 decision outcomes.”
- This is an excellent use of AI.

But the underlying measurements must come from deterministic telemetry.

Governance Health Score

If we use a governance score, Soft Labs should make its calculation visible.

Example:

- **Governance Health**
 - Policy Compliance: 25%
 - Control Effectiveness: 25%
 - Incident Exposure: 20%
 - Trace Completeness: 10%
 - Human Oversight: 10%
 - Data Compliance: 10%
 - Tenant may configure weights.
 - Do not allow an LLM to simply invent:
 - **Governance Score = 94**
 - without a deterministic calculation.
-

Important Boundary — Security Monitoring

Board 4 monitors **AI-specific governance/security risks**.

It should not replace TICVAI's broader:

- Cybersecurity monitoring
- SIEM
- Infrastructure security
- IAM
- Network security
- SOC processes

AI-related security events can integrate with those systems.

Important Boundary — Privacy

Board 4 monitors whether AI data usage complies with configured policies.

It does not replace the entire enterprise privacy/data-governance solution.

Important Boundary — AI Platform Monitoring

AI Platform will later own technical metrics such as:

- Provider Availability
- Token Usage
- Model Endpoint Health

- • Compute
- • Provider Cost
- • Rate Limits
- • Routing

Board 4 only consumes those metrics when they create a **governance or risk impact**.

Example:

Provider latency alone:

AI Platform

Provider failures causing Recommendation AI fallback rate to exceed governance threshold:

Board 4

Important Boundary — Recommendation Analytics

Recommendation Board 5 owns:

- • CTR
- • Conversion
- • Revenue
- • Incrementality
- • Experiments
- • Recommendation Optimization

Board 4 only monitors those metrics when they become governance indicators.

Example:

- Low conversion:
- **Recommendation Board 5**
- Sudden 500% recommendation distribution shift after a model update:

AI Governance Board 4

Board 4 — Final Screen Map

AI Risk, Compliance & Governance Monitoring

1. **AI Governance Monitoring Command Center**
2. **AI Risk Register & Risk Exposure Management**
3. **AI Governance Control Library & Control Effectiveness**
4. **AI Policy Compliance & Violation Monitoring**
5. **AI Data, Privacy & Usage Compliance Monitoring**
6. **AI Quality, Behavior & Governance Drift Monitoring**
7. **AI Governance Alert & Detection Center**
8. **AI Incident & Remediation Management**
9. **AI Compliance, Assurance & Governance Reporting**
10. **AI Governance Review, Action Plan & Continuous Improvement**

AI Governance Module — COMPLETE ARCHITECTURE

With Board 4, the full module is now **4 boards / 40 screens**:

Board 1 — AI Governance & Policy Management

Define **what AI may do**.



Board 2 — AI Approval, Human Oversight & Intervention Management

Define **where humans must control AI**.



Board 3 — AI Explainability, Decision Trace & Audit

Explain and prove **what AI did and why**.



Board 4 — AI Risk, Compliance & Governance Monitoring

Continuously ensure **AI remains under control after deployment**.

Together:

CONTROL → HUMAN OVERSIGHT → EXPLAIN & PROVE → MONITOR & IMPROVE

That gives TICVAI a centralized governance layer that can sit above every current and future AI capability without duplicating the business logic of Ticketing, Pricing, CRM, Recommendation, Fraud, Forecasting, Access Control, or the future AI Platform.